





Integrating Physics in Learning for Model-Based Control

Rolf Findeisen, Maik Pfefferkorn, Sebastian Hirt, Hoang Hai Nguyen, Andreas Höhl

Control and Cyber-Physical Systems Laboratory
Technical University of Darmstadt, Germany



Motivation






Autonomous systems should be cable of

- operating safely and reliably
- reacting to disturbances & adapt to changes in the environment, operation mode or task objectives

Reality: autonomous systems often fail to adapt ► loss of safety and reliability

Fuse control and machine learning to restore adaptability and achieve autonomy

Motivation



Machine Learning

- Significant advancements in recent years
- Resembles human learning by improving from experience / acquired data
- Naturally adapts to changing conditions

Mainly driven by

- Availability of data for training due to digitalization
- (Cloud) computing resources for learning
- Open-source software (PyTorch, ...)

Fuse control and machine learning to restore adaptability and achieve autonomy

- Safety and reliability ► constraint satisfaction (quality, ...)?
- **Exploit prior knowledge (models, physical insight, ...)?**
 - purely data-based models / controllers do not necessarily reflect physical properties
 - decreased control performance
 - reduced explainability

Model-based Control

- Many control methods rely on **models** of the system, **objectives, constraints, and disturbances**
- Models are traditionally obtained from first principles (physical laws), or via identification

Common model-based approaches

- Inversion-based control/feedback linearization
- Optimization based control (LQR, Model Predictive Control, ...)
- Model-based Reinforcement learning
- ...


Advantages:

- Rich literature, sound theory, guarantees by design and validation possible
- Leverages physical insights, interpretability
- Widely used and accepted



Challenges:

- Time consuming modelling, identification
- Difficult to obtain sufficiently good model if the system is complex, changes, ...


Challenge: Machine learning does not guarantee that the model takes physical insights into account



Motivation



$$\begin{aligned}\dot{x} &= \sigma(x - y) \\ \dot{y} &= x(\rho - z) \\ \dot{z} &= xy - \beta z\end{aligned}$$



Physics-informed Machine Learning

- Physical dynamics models are readily available
 - exploit prior (expert) knowledge
- Accelerate simulations/predictions especially for **high-dimensional** and **complex dynamics**

Approaches

- Provide physical insights to ML models via
 - Cost function during training
 - Structurally via model architecture

Physics-informed ML models provide



- Improved interpretability and explainability
- Enhanced extrapolation, robustness, and generalization (e.g., preventing overfitting)
- Better sample-efficiency due to reduced search space

► **Exploit the strengths of existing physics knowledge and machine learning for model-based control**

Physics-Informed Learning for Model-based Control

Motivation

- Fusing control and machine learning to achieve autonomy of systems
- Physics-informed ML: integrate known parts, infer unknown parts from data




Physics-informed machine learning

- Integrating physics in Gaussian process models
- Integrating physics in neural network models

Physics-informed learning for model-based control

- Physics-informed model predictive control




Physics-informed learning: build reliable model integrating known physics and data

Physics-Informed Learning for Model-based Control

Motivation

- Fusing control and machine learning to achieve autonomy of systems
- Physics-informed ML: integrate known parts, infer unknown parts from data

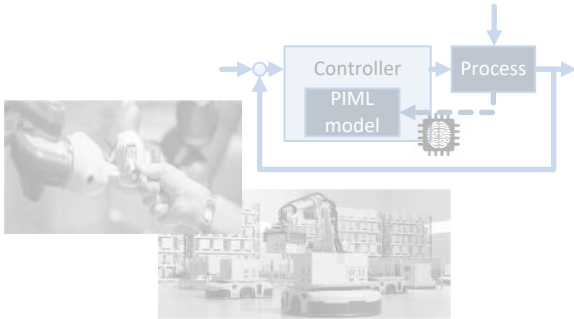


Physics-informed machine learning

- Integrating physics in Gaussian process models
- Integrating physics in neural network models


Physics-informed learning for model-based control

- Physics-informed model predictive control



An Intuition about Gaussian Processes

Core idea: Modeling of realizations of uncertain functions as samples from normal distributions

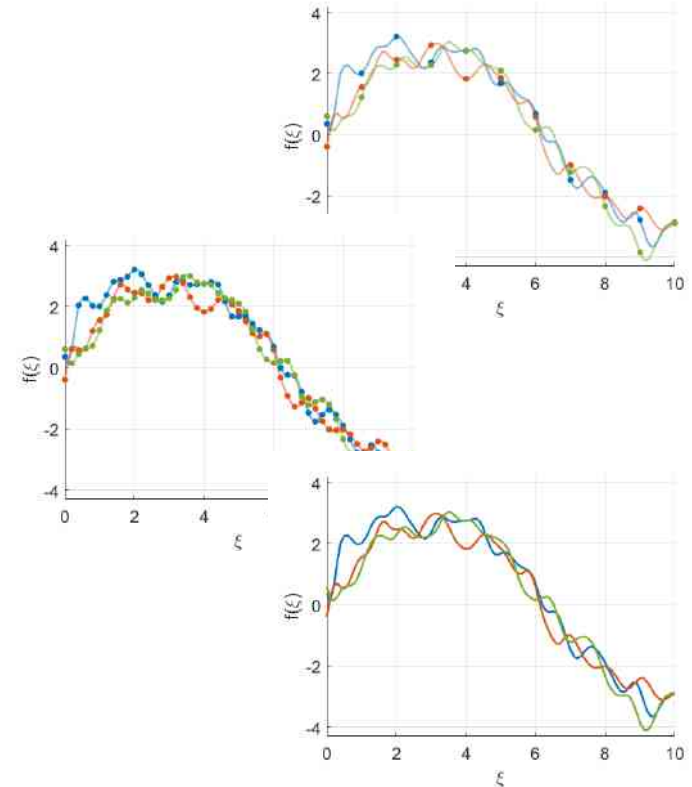


Definition (Gaussian process)

Be T an index set and be Ω a sample space. A Gaussian process is a collection

$$X = \{X_t(\omega) \mid t \in T, \omega \in \Omega\}$$

of random variables $X_t(\omega)$ any finite subset of which has a joint and consistent Gaussian distribution.



Gaussian Process Priors

A Gaussian process $f(\xi) \sim \mathcal{GP}(m(\xi), k(\xi, \xi'))$ is a Gaussian probability distribution over functions

Objective:

- Exploit knowledge provided by **training data** $\{(\xi_i, \gamma_i)_{i=1}^n | \gamma_i = f(\xi_i) + \varepsilon, \varepsilon \sim \mathcal{N}(0, \sigma_n^2)\}$
- Infer unobserved function values (**test targets**) $f_* = f(\xi_*)$ at **test inputs** ξ_*

Basic assumption (joint prior distribution)

Any finite number of function values have a joint Gaussian probability distribution, especially

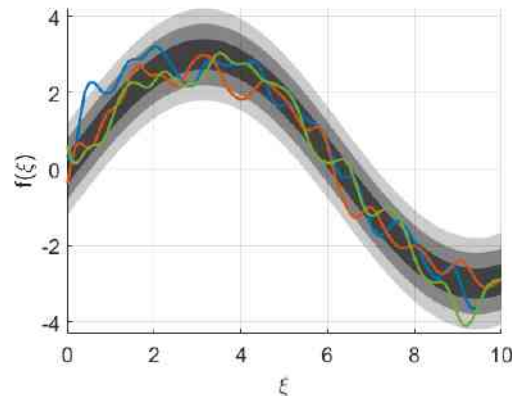
$$\begin{bmatrix} g(\Xi) \\ g(\Xi_*) \end{bmatrix} \sim \mathcal{N} \left(\begin{bmatrix} m(\Xi) \\ m(\Xi_*) \end{bmatrix}, \begin{bmatrix} k(\Xi, \Xi) + \sigma_n^2 \mathbb{I} & k(\Xi, \Xi_*) \\ k(\Xi_*, \Xi) & k(\Xi_*, \Xi_*) \end{bmatrix} \right)$$

where $\Xi = [\xi^{(1)}, \dots, \xi^{(n)}]^\top \in \mathbb{R}^{n \times d}$, $\Xi_* = [\xi_*^{(1)}, \dots, \xi_*^{(n_*)}]^\top \in \mathbb{R}^{n_* \times d}$

Prior GP model of the unknown function

$$g(\xi_*) \sim \mathcal{GP}(m(\xi_*), k(\xi_*, \xi_*'))$$

- Encodes beliefs about the unknown function



Prior Model - Mean Functions


Prior mean function $m : \mathbb{R}^d \rightarrow \mathbb{R}, \xi \mapsto E[f(\xi)]$

$$g(\xi) \sim \mathcal{GP}(m(\xi), k(\xi, \xi'))$$

Used to encode prior knowledge about the underlying function


Constant mean function:

$$m(\xi) = c, c \in \mathbb{R}$$




Affine mean function:

$$m(\xi) = a\xi + b, a, b \in \mathbb{R}$$



Quadratic mean function:

$$m(\xi) = a\xi^2, a \in \mathbb{R}$$



The prior mean functions specifies an overall trend but not the particular shape of the functions

Prior Model: Covariance Functions


Prior covariance (kernel) function $k : \mathbb{R}^d \times \mathbb{R}^d \rightarrow \mathbb{R}, (\xi, \xi') \mapsto \text{Cov}[f(\xi), f(\xi')]$

$$f(\xi) \sim \mathcal{GP}(m(\xi), k(\xi, \xi'))$$

Used to encode prior knowledge about the underlying function

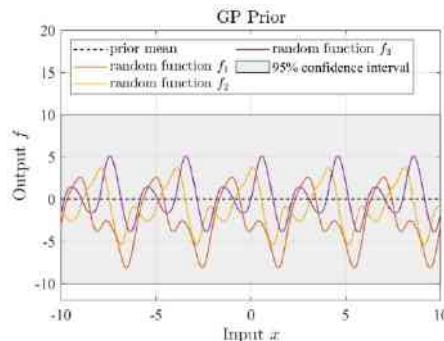
Squared exponential kernel:

► Cont. differentiable fcn's



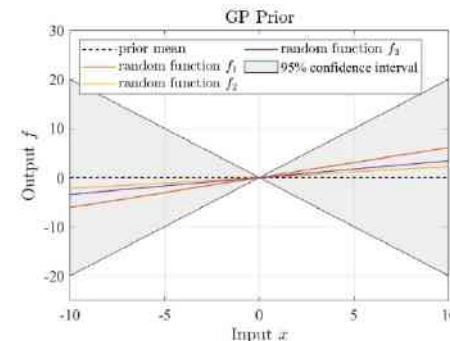
Periodic kernel

► Periodic functions



Linear kernel:

► Linear functions



The prior covariance function allows to induce properties such as smoothness, periodicity, ... of the sample functions and defines the function class to which they belong

Gaussian Process Posteriors

A Gaussian process $f(\xi) \sim \mathcal{GP}(m(\xi), k(\xi, \xi'))$ is a Gaussian probability distribution over functions

Incorporation of training data:

- **Conditioning of the prior on the training data**
 - posterior model / predictive distribution

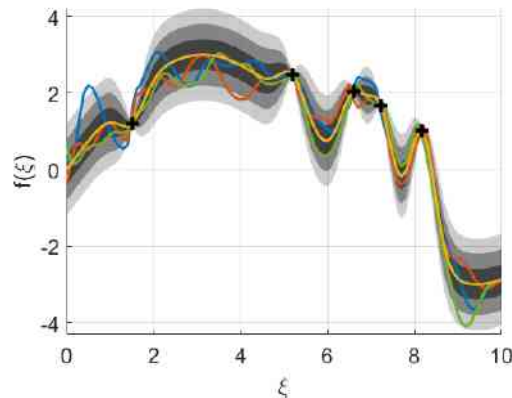
$$g(\Xi_*) \mid \Xi_*, \Xi, \gamma \sim \mathcal{N}(m^+(\Xi_*), k^+(\Xi_*, \Xi_*))$$

- $m^+(\xi_*) = m(\xi_*) + k(\xi_*, \Xi) \Sigma (\gamma - m(\Xi))$
 - Estimate for the unknown function values
- $k^+(\xi_*, \xi'_*) = k(\xi_*, \xi'_*) - k(\xi_*, \Xi) \Sigma k(\Xi, \xi'_*) + \sigma_n^2$
 - $k^+(\xi_*, \xi_*)$ quantifies prediction uncertainty
- $\Sigma = [k(\Xi, \Xi) + \sigma_n^2 \mathbb{I}]^{-1}$

Posterior GP model of the unknown function

$$g(\xi_*) \sim \mathcal{GP}(m^+(\xi_*), k^+(\xi_*, \xi'_*))$$

- **Rejection of prior functions that are inconsistent with the training observations**




Gaussian processes model unknown functions probabilistically based on previous observations

Hyperparameter Learning

- Prior mean $m(\cdot; \theta)$ and covariance $k(\cdot, \cdot; \theta)$ define GP model
- Dependent on **hyperparameters** θ

► **Meaningful prediction model requires suitable hyperparameters**



Learning optimal hyperparameters by maximizing data evidence (logarithmic marginal likelihood):

$$\theta^* = \arg \max_{\theta} \underbrace{\left\{ \frac{1}{2} (\gamma - m(\Xi))^T [k(\Xi, \Xi) + \sigma_n^2 \mathbb{I}]^{-1} (\gamma - m(\Xi)) + \frac{1}{2} \log (\det(k(\Xi, \Xi) + \sigma_n^2 \mathbb{I})) + \frac{n}{2} \log(2\pi) \right\}}_{\log(p(\gamma|\Xi, \theta))}$$

- Corresponds to maximizing the GP's capability of explaining the training data
- Global optimum not guaranteed ► local optima correspond to different interpretations of the data


- **GP learning is an optimization problem**
- **Hyperparameter learning adapts the GP model to the underlying problem**

Physics-Informed Gaussian Processes

How to embed physical prior knowledge in Gaussian process models?

Systems are often partially unknown

- Some **physics-based model** $\bar{\gamma} = \bar{\varphi}(\xi)$ of the unknown function $\gamma = \varphi(\xi)$ **is available**
- This **model is uncertain/imperfect** (e.g., unknown parameters, unmodeled parts, ...)
- The **missing information should be provided by data**



Relying purely on the physics-based model results in an insufficient prediction quality

Relying on a GP model based on only the data may lead to models that are inconsistent with the known physics ► unrealistic predictions

We want to design a GP model that

- **complies with the known physics**, i.e., incorporates the known model
- **corrects/infers the uncertain parts of the model using the data**
- **provides uncertainty information about the model**

Physics-Informed Gaussian Processes

How to embed physical prior knowledge in Gaussian process models?

Physical knowledge may be available in various forms

- Bound constraints
- Monotonicity constraints
- Differential equation constraints
- Boundary condition constraints
- ...

Strategies to bound constraints

- Transforming the GP output
- Employing non-Gaussian likelihoods
- Truncation of the Gaussian output distribution
- Constrained MLE optimization (training)

Strategies to monotonicity constraints

- Constrained likelihood including derivative information
- Truncated Gaussian distributions

Strategies to differential equation constraints

- Block covariance kernels to enforce compliance with DEs at collocation points
- Specialized kernel construction (restriction of the GP sample space)
- Empirical mean and covariance

Strategies to boundary condition constraints

- Spectral expansion of kernels

In the following: focus on **differential equation constraints**

Physics-Informed Gaussian Processes


How to embed physical prior knowledge in Gaussian process models?

Reminder: GP posterior mean

$$m^+(\xi_*) = m(\xi_*) + \sum_{i=1}^n \alpha_i k(\xi, \xi_*)$$

► linear **combination of prior covariance functions “shifted” by prior mean function**

► **Use a dedicated prior to incorporate physics**



Physics-informed prior mean

- Use $m(\xi_*) = \bar{\varphi}(\xi_*)$
- The GP learns the difference between $\bar{\varphi}(\xi_*)$ and $\varphi(\xi_*)$ from the data
- The functional structure of the posterior mean and the sample functions does not match that of the true function

Physics-informed prior covariance

- If the functional structure of φ is known, construct $k(\cdot, \cdot)$ such that $m^+(\cdot)$ and sample functions are restricted to the same class
- Any sample complies with physics, i.e., has an exact real-world realization

How to construct physics-informed GP priors?

Physics-Informed Gaussian Processes

Purely data-driven priors

[X. Yang, D. Barajas-Solano, G. Tartakovsky, A. M. Tartakovsky, 2018]

Stochastic physics-based models:

include expert knowledge (known physics) and random parameters or random fields (imperfect system knowledge)

- stochastic physics-based model $\bar{\varphi}(\xi; \omega)$
- M realizations $\{\Phi^m(\xi)\}_{m=1}^M$ thereof (**Monte-Carlo simulations/sampling**)
- ▶ $\Phi(\xi) \sim \mathcal{GP}(m_{\text{MC}}(\xi), k_{\text{MC}}(\xi, \xi'))$

Construction of a purely data-based prior:

$$m_{\text{MC}}(\xi) = \frac{1}{M} \sum_{m=1}^M \Phi^m(\xi)$$

$$k_{\text{MC}}(\xi, \xi') = \frac{1}{M-1} \sum_{m=1}^M (\Phi^m(\xi) - m_{\text{MC}}(\xi)) \cdot (\Phi^m(\xi') - m_{\text{MC}}(\xi'))$$

GP posterior

$$\begin{aligned} m^+(\xi_*) &= m_{\text{MC}}(\xi_*) \\ &\quad + k_{\text{MC}}(\xi_*, \Xi) k_{\text{MC}}(\Xi, \Xi)^{-1} (\gamma - m_{\text{MC}}(\Xi)) \\ k^+(\xi_*, \xi'_*) &= k_{\text{MC}}(\xi_*, \xi'_*) \\ &\quad - k_{\text{MC}}(\xi_*, \Xi) k_{\text{MC}}(\Xi, \Xi)^{-1} k_{\text{MC}}(\Xi, \xi'_*) \end{aligned}$$

- ▶ All covariances and means are computed empirically from the ensemble

- No need to assume specific structural knowledge or a specific, parametrized prior
- No restriction to certain function classes
- The model fulfills linear physical constraints

The data needs to contain all relevant information

Physics-Informed Gaussian Processes

Embedding differential equations (DE)

[Long et al., 2022]

Example:

$$\partial_t \varphi - \nu \partial_\xi^2 \varphi + \beta \varphi (\varphi^2 - 1) + g(\xi) = 0$$

with target function $\varphi(\xi(t))$

Construct GP priors over φ and g :

$$\varphi(\xi) \sim \mathcal{GP}(m_\varphi(\xi), k_\varphi(\xi, \xi'))$$

$$g(\xi) \sim \mathcal{GP}(m_g(\xi), k_g(\xi, \xi'))$$

Learned functions are supposed to not only fit observations $\{\Xi, \varphi(\Xi)\}$, $\{\Xi, g(\Xi)\}$ but also to comply with the known differential eq.

► **Idea: enforce additionally satisfaction of the L.H.S. of the DE at collocation points $\hat{\Xi}$**

1. Embedding the DE requires derivatives of φ

► GPs are closed under differentiation: derivative models are again GPs

2. Joint Gaussian prior $p(f) = \mathcal{N}(f \mid \mu, \Sigma)$ over $f = [\varphi; \hat{\varphi}; \hat{\varphi}_t; \hat{\varphi}_{\xi\xi}; \hat{g}]$, $\hat{\varphi}_\cdot = \varphi_\cdot(\hat{\Xi})$

3. Noise model: $p(\gamma \mid f) = \mathcal{N}(\gamma \mid f, \sigma_n^2 \mathbb{I})$

4. Virtual second likelihood to integrate DE:

$$p(0 \mid f) = \mathcal{N}(0 \mid \hat{\varphi}_t - \nu \hat{\varphi}_{\xi\xi} + \beta \hat{\varphi} \circ (\hat{\varphi} \circ \hat{\varphi} - 1) + g, \lambda \mathbb{I})$$

► **Joint prior model:**

$$p(f, \gamma, 0) = \mathcal{N}(f \mid \mu, \Sigma) \mathcal{N}(\gamma \mid f, \sigma_n^2 \mathbb{I}) \cdot \mathcal{N}(0 \mid \hat{\varphi}_t - \nu \hat{\varphi}_{\xi\xi} + \beta \hat{\varphi} \circ (\hat{\varphi} \circ \hat{\varphi} - 1) + g, \lambda \mathbb{I})$$

Tune λ to enforce conformity to physics

For dedicated functional structures, DE embeddings can be done by design instead of a penalty

Physics-Informed Gaussian Processes

Physically consistent GP models of conservative Lagrangian systems

[Evangelisti, Hirche, 2022]

Euler-Lagrange Dynamic Systems

$$\frac{d}{dt}(\nabla_{\dot{q}}L) - \nabla_q L = \tau$$

with Lagrangian $L = T(q, \dot{q}) - V(q)$

- Kinetic energy: $T(q, \dot{q}) = \dot{q}^\top M(q) \dot{q}$
- Potential energy: $V(q) = U(q) + G(q)$
- Elastic energy: $U(q) = q^\top S(q) q$
- Gravitational energy: $G(0) = 0, \nabla_q G(0) = 0$

► Alternatively represented by

$$\tau = M(q)\ddot{q} + C(q, \dot{q})\dot{q} + g(q)$$

Lagrange differential operator

$$\mathcal{L}_q = \frac{d}{dt} \nabla_{\dot{q}} - \nabla_q = (\nabla_{\dot{q}}^\top \ddot{q} + \nabla_q^\top \dot{q}) \nabla_{\dot{q}} - \nabla_q$$

Model kinetic, elastic and gravitational energies by Gaussian processes

$$T(q, \dot{q}) \sim \mathcal{GP}(m_T(q, \dot{q}), k_T(q, q', \dot{q}, \dot{q}'))$$

$$U(q) \sim \mathcal{GP}(m_U(q), k_U(q, q'))$$

$$G(q) \sim \mathcal{GP}(m_G(q), k_G(q, q'))$$

GPs are closed under linear operators, yielding

$$\tau \sim \mathcal{GP}(m_\tau(q, \dot{q}, \ddot{q}), k_\tau(q, q', \dot{q}, \dot{q}', \ddot{q}, \ddot{q}'))$$

$$m_\tau = \mathcal{L}_q m_L, \quad k_\tau = \mathcal{L}_q \mathcal{L}_{q'}^\top k_L$$

$$m_L = m_T - m_V, \quad k_L = k_T + k_V$$

$$m_V = m_U + m_G, \quad k_V = k_U + k_G$$

Physics-Informed Gaussian Processes

Physically consistent GP models of conservative Lagrangian systems

[Evangelisti, Hirche, 2022]

Properties of Euler-Lagrange Dynamic Systems

1. Kinetic energy: $T(q, \dot{q}) = \dot{q}^\top M(q) \dot{q}$
 $M(q) = M^\top(q)$, $M(q) \succ 0$
 2. Elastic energy: $U(q) = q^\top S(q) q$
 $S(q) = S^\top(q)$, $S(q) \succ 0$
 3. Gravitational energy: $G(0) = 0$, $\nabla_q G(0) = 0$
- How to enforce compliance of the GP model with the known system properties?

Define GP prior over gravitational energy such that $m_G(0) = 0$, $\nabla_q m_G(0) = 0$

- Compliance with the equilibrium conditions on the gravitational energy
- can be easily ensured

Prior mean for kinetic and elastic energies:

$$m_T(q, \dot{q}) = \frac{1}{2} \dot{q}^\top M_0(q) \dot{q}, \quad m_U(q) = \frac{1}{2} q^\top S_0(q) q$$

Prior covariances for kinetic and elastic energies

$$k_T(q, q', \dot{q}, \dot{q}') = \frac{1}{4} (\dot{q} \circ \dot{q})^\top \Theta_T(q, q') (\dot{q} \circ \dot{q}')$$

$$k_U(q, q') = \frac{1}{4} (q \circ q)^\top \Theta_U(q, q') (q \circ q')$$

$$\Theta_{\cdot}(q, q') = R_{\cdot}^\top(q, q') R_{\cdot}(q, q')$$

- $R_{\cdot}(q, q')$ is an upper triangular matrix of kernels

$$r(q, q') = \sigma_f^2 \exp \left(-\frac{1}{2} (q - q')^\top \Lambda^{-1} (q - q') \right)$$

The physically consistent GP possesses the properties of EL systems (with high probability)

Physics-Informed Gaussian Processes

Physically consistent GP models of port-Hamiltonian systems

[Beckers et al., 2023]

Port-Hamiltonian systems:

$$\dot{x} = [J(x) - R(x)]\nabla_x H(x) + G(x)u$$

$$y = G(x)^\top \nabla_x H(x)$$

- Hamiltonian $H(x)$ represents total energy
- Interconnection matrix $J(x)$
- Dissipation matrix $R(x)$
- $G(x)$ specifies interactions with environment
- Port-Hamiltonian formalism for energy-based modeling of multi-physics systems with energy exchange among different domains

The specific structure of port-Hamiltonian systems can be integrated in GP-based models

Port-Hamiltonian GP prior

- Place a GP prior over the Hamiltonian

$$\hat{H} \sim \mathcal{GP}(0, k(x, x'))$$

with SE kernel $k(x, x') = \sigma_f^2 \exp(-\|x - x'\|_{\Lambda^{-1}}^2)$

- Systems dynamics are an affine transformation of the Hamiltonian:

$$\dot{x} \sim \mathcal{GP}(G(x)u, k_{\text{PHS}}(x, x'))$$

$$k_{\text{PHS}}(x, x') = \sigma_f^2 [J(x) - R(x)]$$

$$\cdot \nabla_x^2 k(x, x') [J(x) - R(x)]^\top$$


This restricts the sample space of the GP to port-Hamiltonian systems ► any realization drawn from the GP is guaranteed to represent a PHS

The physics-informed GP model is guaranteed to have the same system properties almost surely

Physics-Informed Learning for Model-based Control

Motivation

- Fusing control and machine learning to achieve autonomy of systems
- Physics-informed ML: integrate known parts, infer unknown parts from data




Physics-informed machine learning

- Integrating physics in Gaussian process models
- Integrating physics in neural network models

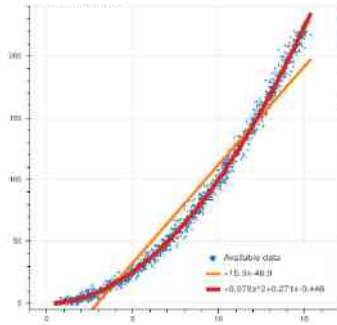
Physics-informed learning for model-based control

- Why model-based control and why physics-informed models?
- Physics-informed model predictive control




Neural Networks Refresher

Linear and nonlinear least-squares regression



- Linear least-squares regression given a set of data: $\gamma = a\xi + b$
- Extension of this set-up to nonlinear functions:
$$\gamma = \sum_{i=0}^n c_i \phi_i(\xi)$$
- **Problem:** The correct parametric model must be available

Neural networks




- **Idea: Use repeated regression to introduce required nonlinearities**
- Input processing by passing it subsequently through several individual regression problems
- Application in Robotics, Automation, ...
- Bio-inspired model architecture

Neural networks are nonparametric, universal function approximators

Neural Networks: Basic Idea

General architecture of (feedforward) neural networks



Feedforward NNs: only forward connections between neurons

- Input layer: incoming values $z^{(0)} = \xi$
- Every **hidden layer neuron** in the network is of the form

$$z_j^{(i)} = \sigma \left(A_j^{(i)} z^{(i-1)} + B_j^{(i)} \right)$$

i : Number of the hidden layer

j : Number of the neuron


$A^{(i)}$: Parameter matrix of layer i

$B^{(i)}$: Bias vector of layer i

- Activation function $\sigma(\cdot)$ has often some switch-like characteristics
- Characterization of different architectures by types of connections and number of layers

Neural Networks: Basic Idea

General architecture of (feedforward) neural networks



Objective:

- Employ the neural network $\hat{\gamma} = \text{NN}(\xi)$ to learn a representation of an unknown function
$$\gamma = \varphi(\xi)$$
- ▶ Infer unobserved function values $\gamma_* = \varphi(\xi_*)$ at test inputs ξ_*


- **The NN outputs depend on the parameters**, i.e., $\hat{\gamma} = \text{NN}(\xi; \theta)$ with parameters
$$\theta = \{(A^{(i)}, B^{(i)}) \mid i = 1, \dots, n_l\}$$
- ▶ The parameters need to be adapted to match the underlying problem
- ▶ Train the network on available data

Feedforward NNs: only forward connections between neurons

How to train the parameters of the network such that it learns to represent the underlying function?

Neural Networks: Training

How can we determine suitable network parameters from training data?



Given training data $\{(\xi_i, \gamma_i) \mid i = 1, \dots, n\}$ and Network architecture $\text{NN}(\cdot; \theta)$, iterate:

- i. **Propagation** of the training inputs through the network: $\hat{\gamma}_i = \text{NN}(\xi_i; \theta)$, $i = 1, \dots, n$
- ii. **Compute the loss, i.e., a measure for the NN's error**, e.g.: $L = \frac{1}{n} \sum_{i=1}^n \|\gamma_i - \hat{\gamma}_i\|_2^2$
- iii. **Backpropagation** of the loss through the network layer by layer to compute the gradients $\nabla_{\theta} L$
- iv. **Parameter update** by stochastic gradient descent algorithm exploiting $\nabla_{\theta} L$

Training a neural network is an optimization problem. Similarly to GPs, one can

- embed physical knowledge in the loss function / optimization criterion
- enforce function classes via network structure
- ...


Constrained Neural Networks

- **Restrict parameters or architecture** such that the NN is, e.g.,
 - Convex
 - Lipschitz continuous
 - Monotonous
 - Differentiable
- **Enforce via**
 - Hard-constraints
 - Via **cost function** / regularization
 - Structure of the network itself

- Allows to incorporate prior knowledge
- Can be used to enforce system properties

Example: Convex NN

$$\text{Fit } y(x) = x^2 + \epsilon, \quad \epsilon \sim \mathcal{N}(\mu, \sigma^2)$$



Physics-informed Neural Networks (PINNs)

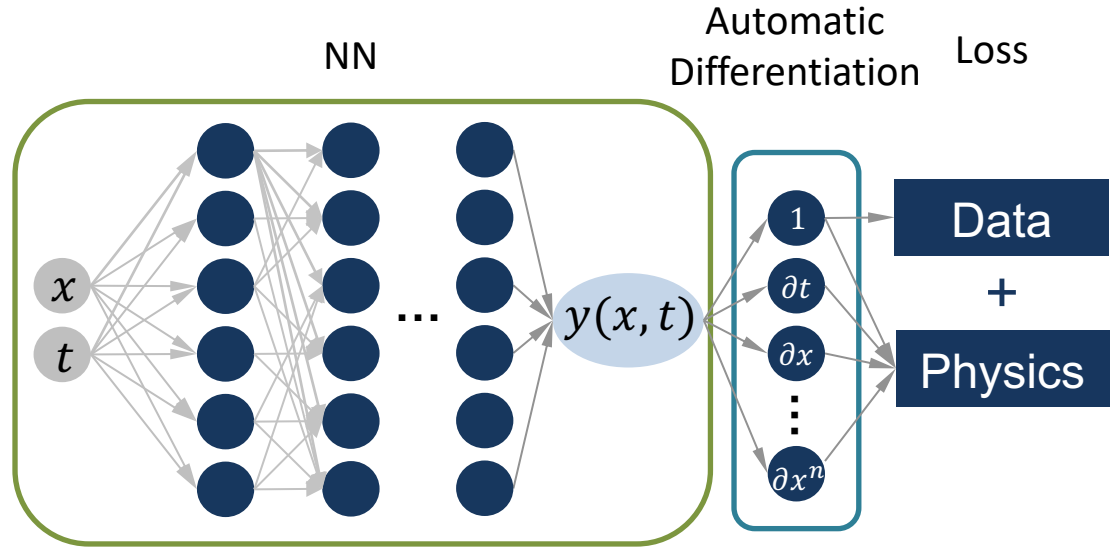
Model needs to satisfy differential operator (physical insight)

$$\mathcal{F}(y(x, t), \frac{\partial y}{\partial t}, \frac{\partial y}{\partial x}, \dots, \frac{\partial^n y}{\partial x^n}) = 0$$

Penalize loss using system model

$$\mathcal{L}(\theta, \mathcal{D}) = \sum_i (y_i - \text{NN}(x_i, t_i; \theta))^2 + \lambda \sum_i \mathcal{F}(\cdot)^2|_{x_i, t_i}$$

- Improves sample efficiency by regularization
- Predictions generally do not strictly comply with physics



Restricted to autonomous systems ► How to extend to controlled systems?

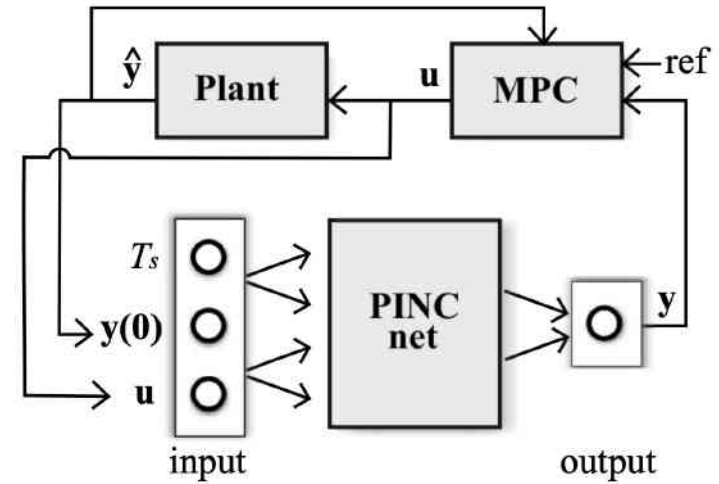
Physics-informed Neural Networks for Control (PINCs)

“PINNs do not allow control inputs, neither can they simulate for variable long-range intervals without serious degradation in their predictions”

- Same basic principle as PINNs
- Introduce control input u
- Train the network for a sampling period T_s
- NN learns the solution of the system dynamics
- Multi-step predictions by repeated evaluations

Related to / motivated by direct multiple shooting

- Can be combined with model-based control
- Allows for long time prediction



[Antonelo et al. 2022]

Example: PINCs for Control of Multi-Link Manipulators

Use PINC to learn forward dynamics of multilink manipulator.

Given the generalized coordinates $q = [\alpha, \beta]^T$ and the input currents u , the system can be described by:

$$0 = \begin{bmatrix} I & 0 \\ 0 & M(q) \end{bmatrix} \begin{bmatrix} \dot{q} \\ \ddot{q} \end{bmatrix} - \begin{bmatrix} \dot{q} \\ h(q, \dot{q}) - k(q, \dot{q}) \end{bmatrix} - \begin{bmatrix} 0 \\ Bu \end{bmatrix}$$

- With $x = [q \quad \dot{q}]^T$, we denote $0 = \mathcal{F}(x, \dot{x}, u)$
- zero-order hold for u and sampling time T_s


► Solution: $x_{k+1} = \phi(T_s, u_k, x_k)$

Learn DNN to approximate $\hat{x}(t) = \hat{\phi}_\theta(t, u_k, x_0) \approx \phi(t, u_k, x_0)$

Sample data $x_i = \phi(t_i, u_i, x_{0,i})$ and select weights θ to minimize

$$\sum_i \left\| \hat{\phi}_\theta(t_i, u_i, x_{0,i}) - x_i \right\|^2 + \lambda \left\| \mathcal{F}(\hat{\phi}_\theta(t_i, u_i, x_{0,i}), \frac{\partial}{\partial t} \hat{\phi}_\theta(t_i, u_i, x_{0,i}), u_i) \right\|^2$$

PINC can achieve good data fit and conformity with physics, traded off by λ







[Nicodemus et al. 2022]

Example: PINCs for Control of Multi-Link Manipulators

Evaluate the model and combine with MPC to track a reference trajectory.

- Evaluation of PINN is faster than solving the dynamics
- Enables for example real time MPC




[Nicodemus et al. 2022]

PINCs enable the efficient combination of prior system knowledge and data for control tasks

Physics-Informed Learning for Model-based Control

Motivation

- Fusing control and machine learning to achieve autonomy of systems
- Physics-informed ML: integrate known parts, infer unknown parts from data




Physics-informed machine learning

- Integrating physics in Gaussian process models
- Integrating physics in neural network models


Physics-informed learning for model-based control

- Physics-informed model predictive control (same examples)




Fusing Learning and Model-based Control

Model-based control




- + Exploits physical knowledge
- + Provides physical insight
- + Exploits model and data
- + Guarantees possible
- difficult to handle large changes

Learning-supported control



- + Exploits physical knowledge
- + Provides physical insight
- + Exploits model and data
- + Guarantees possible
- + Allows adaptation


Learning-based control




- + Exploits available data
- + No physical model needed
- Lots of data needed
- Physical insight, guarantees?
- Retraining required

The performance of controllers can be significantly improved if physics-informed models are employed


References & Disturbances




Model



Cost & Constraints



Controller




Model-based Control: Model Predictive Control

MPC repeatedly solves a finite-horizon optimal control problem:

$$\begin{aligned} \min_{u_k} \quad & \left\{ \sum_{i=0}^{N-1} F(x_{i|k}, u_{i|k}) + V_f(x_{N|k}) \right\} \\ \text{s. t. } \quad & \forall i \in \mathcal{I}_{0:N-1} : x_{i+1|k} = f(x_{i|k}, u_{i|k}), \quad x_{0|k} = x_k \\ & \forall i \in \mathcal{I}_{0:N-1} : x_{i|k} \in \mathcal{X}, \quad u_{i|k} \in \mathcal{U} \\ & x_{N|k} \in \mathcal{E} \end{aligned}$$

1) Obtain x_k & solve OCP at time point t_k




Model-based Control: Model Predictive Control

MPC means to repeatedly solve a finite-horizon optimal control problem:

$$\begin{aligned} \min_{\mathbf{u}_k} \quad & \left\{ \sum_{i=0}^{N-1} F(x_{i|k}, u_{i|k}) + V_f(x_{N|k}) \right\} \\ \text{s. t. } \quad & \forall i \in \mathcal{I}_{0:N-1} : x_{i+1|k} = f(x_{i|k}, u_{i|k}), \quad x_{0|k} = x_k \\ & \forall i \in \mathcal{I}_{0:N-1} : x_{i|k} \in \mathcal{X}, \quad u_{i|k} \in \mathcal{U} \\ & x_{N|k} \in \mathcal{E} \end{aligned}$$

1) Obtain x_k & solve OCP at time point t_k

2) Employ $u(t) = u_{0|k}$ for $t \in [t_k, t_{k+1}]$



Model-based Control: Model Predictive Control


MPC means to repeatedly solve a finite-horizon optimal control problem:

$$\begin{aligned} \min_{u_k} \quad & \left\{ \sum_{i=0}^{N-1} F(x_{i|k}, u_{i|k}) + V_f(x_{N|k}) \right\} \\ \text{s. t. } \quad & \forall i \in \mathcal{I}_{0:N-1} : x_{i+1|k} = f(x_{i|k}, u_{i|k}), \quad x_{0|k} = x_k \\ & \forall i \in \mathcal{I}_{0:N-1} : x_{i|k} \in \mathcal{X}, \quad u_{i|k} \in \mathcal{U} \\ & x_{N|k} \in \mathcal{E} \end{aligned}$$

1) Obtain x_k & solve OCP at time point t_k

2) Employ $u(t) = u_{0|k}$ for $t \in [t_k, t_{k+1}]$

3) Shift horizon by one step, i.e., $t_k \leftarrow t_{k+1}$




Model-based Control: Model Predictive Control

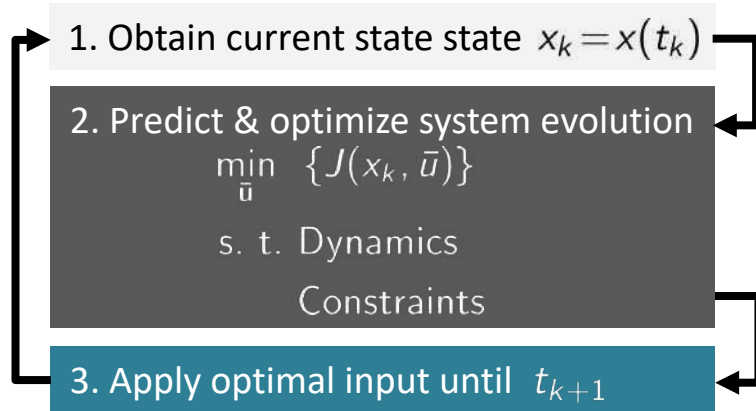
MPC means to repeatedly solve a finite-horizon optimal control problem:

$$\begin{aligned} \min_{u_k} \quad & \left\{ \sum_{i=0}^{N-1} F(x_{i|k}, u_{i|k}) + V_f(x_{N|k}) \right\} \\ \text{s. t. } \quad & \forall i \in \mathcal{I}_{0:N-1} : x_{i+1|k} = f(x_{i|k}, u_{i|k}), \quad x_{0|k} = x_k \\ & \forall i \in \mathcal{I}_{0:N-1} : x_{i|k} \in \mathcal{X}, \quad u_{i|k} \in \mathcal{U} \\ & x_{N|k} \in \mathcal{E} \end{aligned}$$

- 1) Obtain x_k & solve OCP at time point t_k
- 2) Employ $u(t) = u_{0|k}$ for $t \in [t_k, t_{k+1}]$
- 3) Shift horizon by one step, i.e., $t_k \leftarrow t_{k+1}$



Model-based Control: Model Predictive Control



$$\min_{\bar{u}_k} \left\{ \int_{t_k}^{t_k+T} F(\tau, \bar{x}(\tau), \bar{u}(\tau)) + E(\bar{x}(\tau + T)) \right\}$$


$$\text{s. t. } \forall \tau \in [t_k, t_k + T] :$$

$$\dot{\bar{x}}(\tau) = f(\bar{x}(\tau), \bar{u}(\tau)) , \quad \bar{x}(t_k) = x(t_k)$$

$$\bar{y}(\tau) = h(\bar{x}(\tau))$$

$$\bar{u}(\tau) \in \mathcal{U}, \bar{x}(\tau) \in \mathcal{X}$$

$$\bar{x}(t_k + T) \in \Omega$$



Advantages of MPC

- Direct consideration of constraints ► **safety**
- Nonlinear, multiple inputs/output systems ► **flexibility**
- Use of preview information and adaptation of model/parameters ► **adaptiveness**
- Possibility to “optimize” a cost ► **performance**
- Many stability & robustness results available ► **stability & robustness**
- **Efficient embedded optimization strategies**




Perfect knowledge of system & environment required ► use machine learning to adapt to the unknown

Example: “Physics”-Informed MPC

Objective: increased autonomy by learning from interaction

► Learn **reference signal** or **disturbance model** to support desired interaction with environment

Examples



www.interactanalysis.com/

Control Task (Reference Trajectory Tracking)

Given output reference $r_{tt} : [0, \infty) \rightarrow \mathbb{R}^{n_y}$, $t \mapsto y^r(t)$ and system $\dot{x}(t) = f(x(t), u(t))$, $y(t) = h(x(t))$, design controller that achieves

i. Convergence to reference:

$$\lim_{t \rightarrow \infty} \|y(t) - r_{tt}(t)\| = 0$$

ii. Constraint satisfaction:

$$\forall t \geq 0 : x(t) \in \mathcal{X}, u(t) \in \mathcal{U}, y(t) \in \mathcal{Y}$$

The reference is not always known

Support controller with a learned reference to ensure safe and desired interaction with environment

Example: Physics-Informed MPC

Learning Task: Reference Generator for Trajectory Tracking

Learn a reference generator $g : [t_k, t_k + T] \rightarrow \mathcal{Y} \subset \mathbb{R}^{n_y}$ from data \mathcal{D} , which provides the reference $y^r(\tau) = g(\tau; \mathcal{D})$, $\forall \tau \in [t_k, t_k + T]$ that is used in the MPC, such that

i. **Trackability:** The reference is trackable under output, state and input constraints, i.e.,



$$\forall \tau \in [t_k, t_k + T] : \exists u^r(\tau) \in \mathcal{U} : y^r(\tau) = h(x^r(\tau)) \in \mathcal{Y}, x^r(\tau) \in \mathcal{X}$$

ii. **Data fit:** The reference is data-consistent (and preferably of low complexity)

$$\forall i = 1, \dots, N : y^r(t_i) \approx r_i, \text{ where } (t_i, r_i) \in \mathcal{D}$$

Data fit: Use GP as reference generator

Trackability & constraint satisfaction: Constrain learning process exploiting system properties




Data fit is trivially achieved; How to ensure trackability / safety of the learned reference?

Example: Physics-Informed MPC

Assumption: the system under consideration is differentially flat

$$\exists \Phi, \Psi : x_r = \Phi \left(y_r, \dot{y}_r, \dots, y_r^{(\beta-1)} \right)$$

$$u^r = \Psi \left(y_r, \dot{y}_r, \dots, y_r^{(\beta)} \right)$$



Integrating physics by constrained hyperparameter optimization

$$\theta^* = \arg \max_{\theta} \{ \log(p(\gamma \mid \chi, \theta)) \}$$

s. t. $\forall \tau \in [0, \bar{t}] :$

$$m^+(\tau; \mathcal{D}, \theta) \in \mathcal{Y}$$

$$\Phi \left(m^+(\tau; \mathcal{D}, \theta), \dot{m}^+(\tau; \mathcal{D}, \theta), \dots, m^{+(\beta-1)}(\tau; \mathcal{D}, \theta) \right) \in \mathcal{X}$$

$$\Psi \left(m^+(\tau; \mathcal{D}, \theta), \dot{m}^+(\tau; \mathcal{D}, \theta), \dots, m^{+(\beta)}(\tau; \mathcal{D}, \theta) \right) \in \mathcal{U}$$

Standard training

System dynamics

Constraints

Theorem [Matschek *et al.*, 2021]:

Feasibility of hyperparameter optimization guarantees trackability and constraint satisfaction for

$[0, T]$

Example: Physics-Informed MPC

Can we obtain a safe reference trajectory (constraint satisfaction) all times?

Example periodic references ► use stationary, periodic kernel $k(t, t') = k(t, t' + nT_p)$

Constrained hyperparameter optimization

$$\theta^* = \arg \max_{\theta} \{ \log(p(\gamma \mid \chi, \theta)) \}$$

s. t. $\forall \tau \in [0, \bar{t}] :$

$$m^+(\tau; \mathcal{D}, \theta) \in \mathcal{Y}$$

$$\Phi \left(m^+(\tau; \mathcal{D}, \theta), \dot{m}^+(\tau; \mathcal{D}, \theta), \dots, m^{+(\beta-1)}(\tau; \mathcal{D}, \theta) \right) \in \mathcal{X}$$

$$\Psi \left(m^+(\tau; \mathcal{D}, \theta), \dot{m}^+(\tau; \mathcal{D}, \theta), \dots, m^{+(\beta)}(\tau; \mathcal{D}, \theta) \right) \in \mathcal{U}$$

$$T_p \leq \bar{t}$$

Coverage constraint

- Ensure safety & trackability for at least one period
- Safety & trackability for all periods

Trackability always ► recursive feasibility & stability certificates of MPC via standard approaches
[Matschek *et al.*, 2021]

Example: Physics-Informed MPC for Mobile Robots

Control Task




Synchronize motion of ego robot to that of target robot, which is unknown to ego

- Learn motion of ego robot as reference

Model

$$\begin{aligned} \dot{x}_1 &= u_1 \sin(x_3) & x_r &= \Phi(y_r, \dot{y}_r) = \left(y_{r,1}, y_{r,2}, \operatorname{atan} \left(\frac{\dot{y}_{r,2}}{\dot{y}_{r,1}} \right) + i\pi \right)^T \\ \dot{x}_2 &= u_1 \cos(x_3) \\ \dot{x}_3 &= u_2 & u_r &= \Psi(y_r, \dot{y}_r, \ddot{y}_r) = \left(\|\dot{y}_r\|_2^2, \frac{\ddot{y}_{r,2}\dot{y}_{r,1} - \dot{y}_{r,2}\ddot{y}_{r,1}}{\|\dot{y}_r\|_2^2} \right)^T \\ y &= (x_1, x_2) \end{aligned}$$


Learned reference trajectories – Outputs and inputs




Unconstrained GP learning yields an unsafe reference trajectory that violates constraints

Example: Physics-Informed MPC for Mobile Robots

State tracking errors







Applied control inputs



Exploiting physically constrained learning for reference generation allows to significantly improve control performance

Example: Physics-Informed MPC for Medical Robotics

Medical Task: Robot-assisted radio-frequency ablation of spinal cord tumors (i.e., heat ablation)



Challenges

- Complete ablation of tumor despite uncertainties
- Avoid thermal overheating of spinal cord and damage to non-tumor-tissue
- Satisfaction of constraints
- “Humans in the loop”

Requirements

- Precise positioning and stabilization of needles for ablation in the spinal cord
- Precise control of applied forces
- Adaptation to changing conditions
- Disturbance rejection/attenuation

Robot supported intervention

Example: Physics-Informed MPC for Medical Robotics



Many subtasks have to be solved

- Planning optimal and collision-free paths
- Selecting/Planning optimal robot configurations
- Tracking of desired needle positions
- Measuring needle positions, registration of events
- Human-Robot-Interface
- Controlling contact forces

Therapy
planning



Path-following
under constraints




Force limitation
&
compensation



Exp. Design
Optimal Control



Constrained
path-following control



Force control



The patient is breathing
► Motion compensation
(unknown reference)



Example: Physics-Informed MPC for Medical Robotics

Learning patient-specific references for breathing motion compensation using GPs


Clinical data

- Time-invariance of hyperparameters
- ▶ Patient-specific offline estimation

Breathing motion modeling

- Individual GP with **periodic kernel** for each quantity; coupling via constraints
- **Smooth** predictions/reference signals
- Extrapolation with decent quality due to selected prior

Improved performance and guaranteed trackability using constrained GP



Use patient-specific hyperparameters and safe data updates for online reference predictions

Example: Physics-Informed MPC for Medical Robotics

Results on motion compensation using GP-supported MPC

MPC

- Horizon length: 30-time steps
- Sampling time: 4 ms
- Reference: GP mean predictions; variances not (directly) exploited


Average control errors

$$e_x = 0.97 \text{ mm}, e_A = 0.23^\circ$$

$$e_y = 0.20 \text{ mm}, e_B = 0.10^\circ$$

$$e_z = 0.78 \text{ mm}, e_C = 0.19^\circ$$

Low tracking errors and constraint satisfaction at all times





Physics-informed learning-supported MPC for motion compensation provides very promising results

Conclusions and Outlook


Physics-informed learning

- Models incorporate physical knowledge to better represent real-world systems
- GPs and neural networks can be enforced to comply with physical knowledge by design or by penalty
- Interpretable high-quality models



Model-based control using physics-informed ML models

- Adaptiveness and uncertainty quantification (ML part) while explainable and reasonable behavior (physics)
- Improved control performance: desired system behavior is physically realizable



Outlook

- Physics-informed learning of closed-loop models (controller + model + constraints + ...) with desirable properties

Integrating Physics in Learning for Model-Based Control

Rolf Findeisen, Maik Pfefferkorn, Sebastian Hirt, Hoang Hai Nguyen, Andreas Höhl

Control and Cyber-Physical Systems Laboratory
Technical University of Darmstadt, Germany

